

1. Python for ML/AI

- 1.1. Why Python?
- 1.2. Setup
 - 1.2.1. Install Python.
 - 1.2.2. Installing packages: numpy, pandas, scipy, matplotlib, seaborn, sklearn)
 - 1.2.3. iPython setup.
- 1.3. Introduction
 - 1.3.1. Keywords and Identifiers
 - 1.3.2. Statements, Indentation and Comments
 - 1.3.3. Variables and Datatypes
 - 1.3.4. Input and Output
 - 1.3.5. Operators
- 1.4. Flow Control
 - 1.4.1. If...else
 - 1.4.2. while loop
 - 1.4.3. for loop
 - 1.4.4. break and continue
- 1.5. Data Structures
 - 1.5.1. Lists
 - 1.5.2. Tuples
 - 1.5.3. Dictionary
 - 1.5.4. Strings
 - 1.5.5. Sets
- 1.6. Functions
 - 1.6.1. Introduction
 - 1.6.2. Types of functions
 - 1.6.3. Function Arguments
 - 1.6.4. Recursive Functions
 - 1.6.5. Lambda Functions
 - 1.6.6. Modules
 - 1.6.7. Packages
- 1.7. File Handling
- 1.8. Exception Handling
- 1.9. Debugging Python
- 1.10. NumPy
 - 1.10.1. Introduction to NumPy.
 - 1.10.2. Numerical operations.
- 1.11. Matplotlib
- 1.12. Pandas
 - 1.12.1. Getting started with pandas
 - 1.12.2. Data Frame Basics
 - 1.12.3. Key Operations on Data Frames.

- 1.13. Computational Complexity: an Introduction
 - 1.13.1. Space and Time Complexity: Find largest number in a list
 - 1.13.2. Binary search
 - 1.13.3. Find elements common in two lists.
 - 1.13.4. Find elements common in two lists using a Hashtable/Dict
 - 1.13.5. Further reading about Computational Complexity [Please add a section with these links for reference]
 - 1.13.5.1. <https://medium.com/omarelgabrys-blog/the-big-scary-o-notation-c9352d827ce>
 - 1.13.5.2. <https://rob-bell.net/2009/06/a-beginners-guide-to-big-o-notation/>
 - 1.13.5.3. <https://medium.freecodecamp.org/time-is-complex-but-priceless-f0abd015063c>

2. Plotting for exploratory data analysis (EDA)

- 2.1. Iris dataset
 - 2.1.1. Data-point, vector, observation
 - 2.1.2. Dataset
 - 2.1.3. Input variables/features/dimensions/independent variable
 - 2.1.4. Output Variable/Class Label/ Response Label/ dependent variable
 - 2.1.5. Objective: Classification.
- 2.2. Scatter-plot: 2D, 3D.
- 2.3. Pair plots.
- 2.4. PDF, CDF, Univariate analysis.
 - 2.4.1. Histogram and PDF
 - 2.4.2. Univariate analysis using PDFs.
 - 2.4.3. Cumulative distribution function (CDF)
- 2.5. Mean , Variance, Std-dev
- 2.6. Median, Percentiles, Quantiles, IQR, MAD and Outliers.
- 2.7. Box-plot with whiskers
- 2.8. Violin plots.
- 2.9. Summarizing plots.
- 2.10. Univariate, Bivariate and Multivariate analysis.
- 2.11. Multivariate probability density, contour plot.
- 2.12. Exercise: Perform EDA on Haberman dataset.

3. Probability and Statistics

- 3.1. Introduction to Probability and Stats
 - 3.1.1. Why learn it?
 - 3.1.2. $P(X=x_1)$, Dice and coin example
 - 3.1.3. Random variables: discrete and continuous.
 - 3.1.4. Outliers (or) extreme points.
 - 3.1.5. Population & Sample.
- 3.2. Gaussian/Normal Distribution
 - 3.2.1. Examples: Heights and weights.
 - 3.2.2. Why learn about distributions.

- 3.2.3. Mu, sigma: Parameters
 - 3.2.4. PDF (iris dataset)
 - 3.2.5. CDF
 - 3.2.6. 1-std-dev, 2-std-dev, 3-std-dev range.
 - 3.2.7. Symmetric distribution, Skewness and Kurtosis
 - 3.2.8. Standard normal variate (z) and standardization.
 - 3.2.9. Kernel density estimation.
 - 3.2.10. Sampling distribution & Central Limit theorem.
 - 3.2.11. Q-Q Plot: Is a given random variable Gaussian distributed?
- 3.3. Uniform Distribution and random number generators
 - 3.3.1. Discrete and Continuous Uniform distributions.
 - 3.3.2. How to randomly sample data points. [UniformDisb.ipynb]
- 3.4. Bernoulli and Binomial distribution
- 3.5. Log-normal and power law distribution:
 - 3.5.1. Log-normal: CDF, PDF, Examples.
 - 3.5.2. Power-law & Pareto distributions: PDF, examples
 - 3.5.3. Converting power law distributions to normal: Box-Cox/Power transform.
- 3.6. Correlation
 - 3.6.1. Co-variance
 - 3.6.2. Pearson Correlation Coefficient
 - 3.6.3. Spearman Rank Correlation Coefficient
 - 3.6.4. Correlation vs Causation
- 3.7. Confidence Intervals
 - 3.7.1. Confidence Interval vs Point estimate.
 - 3.7.2. Computing confidence-interval given a distribution.
 - 3.7.3. For mean of a random variable
 - 3.7.3.1. Known Standard-deviation: using CLT
 - 3.7.3.2. Unknown Standard-deviation: using t-distribution
 - 3.7.4. Confidence Interval using empirical bootstrap [BootstrapCI.ipynb]
- 3.8. Hypothesis testing
 - 3.8.1. Hypothesis Testing methodology, Null-hypothesis, test-statistic, p-value.
 - 3.8.2. Resampling and permutation test.
 - 3.8.3. K-S Test for similarity of two distributions.
 - 3.8.4. Code Snippet [KSTest.ipynb]
- 4. Linear Algebra**
 - 4.1. Why learn it ?
 - 4.2. Fundamentals
 - 4.2.1. Point/Vector (2-D, 3-D, n-D)
 - 4.2.2. Dot product and angle between 2 vectors.
 - 4.2.3. Projection, unit vector
 - 4.2.4. Equation of a line (2-D), plane(3-D) and hyperplane (n-D)
 - 4.2.5. Distance of a point from a plane/hyperplane, half-spaces
 - 4.2.6. Equation of a circle (2-D), sphere (3-D) and hypersphere (n-D)

- 4.2.7. Equation of an ellipse (2-D), ellipsoid (3-D) and hyperellipsoid (n-D)
- 4.2.8. Square, Rectangle, Hyper-cube and Hyper-cuboid..

5. Dimensionality reduction and Visualization:

- 5.1. What is dimensionality reduction?
- 5.2. Data representation and pre-processing
 - 5.2.1. Row vector, Column vector: Iris dataset example.
 - 5.2.2. Represent a dataset: $D = \{x_i, y_i\}$
 - 5.2.3. Represent a dataset as a Matrix.
 - 5.2.4. Data preprocessing: Column Normalization
 - 5.2.5. Mean of a data matrix.
 - 5.2.6. Data preprocessing: Column Standardization
 - 5.2.7. Co-variance of a Data Matrix.
- 5.3. MNIST dataset (784 dimensional)
 - 5.3.1. Explanation of the dataset.
 - 5.3.2. Code to load this dataset.
- 5.4. Principal Component Analysis.
 - 5.4.1. Why learn it.
 - 5.4.2. Geometric intuition.
 - 5.4.3. Mathematical objective function.
 - 5.4.4. Alternative formulation of PCA: distance minimization
 - 5.4.5. Eigenvalues and eigenvectors.
 - 5.4.6. PCA for dimensionality reduction and visualization.
 - 5.4.7. Visualize MNIST dataset.
 - 5.4.8. Limitations of PCA
 - 5.4.9. Code example.
 - 5.4.10. PCA for dimensionality reduction (not-visualization)
- 5.5. T-distributed stochastic neighborhood embedding (t-SNE)
 - 5.5.1. What is t-SNE?
 - 5.5.2. Neighborhood of a point, Embedding.
 - 5.5.3. Geometric intuition.
 - 5.5.4. Crowding problem.
 - 5.5.5. How to apply t-SNE and interpret its output (distill.pub)
 - 5.5.6. t-SNE on MNIST.
 - 5.5.7. Code example.

6. Real world problem: Predict sentiment polarity given product reviews on Amazon.

- 6.1. Exploratory Data Analysis.
 - 6.1.1. Dataset overview: Amazon Fine Food reviews
 - 6.1.2. Data Cleaning: Deduplication.
- 6.2. Featurizations: convert text to numeric vectors.
 - 6.2.1. Why convert text to a vector?
 - 6.2.2. Bag of Words (BoW)
 - 6.2.3. Text Preprocessing: Stemming, Stop-word removal, Tokenization, Lemmatization.

- 6.2.4. uni-gram, bi-gram, n-grams.
- 6.2.5. tf-idf (term frequency- inverse document frequency)
- [7.2.5 a] [New Video] Why use log in IDF?
- 6.2.6. Word2Vec.
- 6.2.7. Avg-Word2Vec, tf-idf weighted Word2Vec
- 6.3. Code samples
 - 6.3.1. Bag of Words.
 - 6.3.2. Text Preprocessing
 - 6.3.3. Bi-Grams and n-grams.
 - 6.3.4. TF-IDF
 - 6.3.5. Word2Vec
 - 6.3.6. Avg-Word2Vec and TFIDF-Word2Vec
- 6.4. Exercise: t-SNE visualization of Amazon reviews with polarity based color-coding

7. Classification and Regression Models: K-Nearest Neighbors

- 7.1. Foundations
 - 7.1.1. How “Classification” works?
 - 7.1.2. Data matrix notation.
 - 7.1.3. Classification vs Regression (examples)
- 7.2. K-Nearest Neighbors
 - 7.2.1. Geometric intuition with a toy example.
 - 7.2.2. Failure cases.
 - 7.2.3. Distance measures: Euclidean(L2) , Manhattan(L1), Minkowski, Hamming
 - 7.2.4. Cosine Distance & Cosine Similarity
 - 7.2.5. How to measure the effectiveness of k-NN?
 - 7.2.6. Simple implementation:
 - 7.2.6.1. Test/Evaluation time and space complexity.
 - 7.2.6.2. Limitations.
 - 7.2.7. Determining the right “k”
 - 7.2.7.1. Decision surface for K-NN as K changes.
 - 7.2.7.2. Overfitting and Underfitting.
 - 7.2.7.3. Need for Cross validation.
 - 7.2.7.4. K-fold cross validation.
 - [NEW]8.2.7.4 a Visualizing train, validation and test datasets
 - 7.2.7.5. How to determine overfitting and underfitting?
 - 7.2.7.6. Time based splitting
 - 7.2.8. k-NN for regression.
 - 7.2.9. Weighted k-NN
 - 7.2.10. Voronoi diagram.
 - 7.2.11. kd-tree based k-NN:
 - 7.2.11.1. Binary search tree
 - 7.2.11.2. How to build a kd-tree.

- 7.2.11.3. Find nearest neighbors using kd-tree
 - 7.2.11.4. Limitations.
 - 7.2.11.5. Extensions.
 - 7.2.12. Locality sensitive Hashing (LSH)
 - 7.2.12.1. Hashing vs LSH.
 - 7.2.12.2. LSH for cosine similarity
 - 7.2.12.3. LSH for euclidean distance.
 - 7.2.13. Probabilistic class label
 - 7.2.14. Code Samples for K-NN
 - 7.2.14.1. Decision boundary. [./knn/knn.ipynb and knn folder]
 - 7.2.14.2. Cross Validation.[./knn/kfold.ipynb and knn folder]
 - 7.2.15. Exercise: Apply k-NN on Amazon reviews dataset.
8. Classification algorithms in various situations:
- 8.1. Introduction
 - 8.2. Imbalanced vs balanced dataset.
 - 8.3. Multi-class classification.
 - 8.4. k-NN, given a distance or similarity matrix
 - 8.5. Train and test set differences.
 - 8.6. Impact of Outliers
 - 8.7. Local Outlier Factor.
 - 8.7.1. Simple solution: mean dist to k-NN.
 - 8.7.2. k-distance (A), $N(A)$
 - 8.7.3. reachability-distance(A, B)
 - 8.7.4. Local-reachability-density(A)
 - 8.7.5. LOF(A)
 - 8.8. Impact of Scale & Column standardization.
 - 8.9. Interpretability
 - 8.10. Feature importance & Forward Feature Selection
 - 8.11. Handling categorical and numerical features.
 - 8.12. Handling missing values by imputation.
 - 8.13. Curse of dimensionality.
 - 8.14. Bias-Variance tradeoff.
 - 9.14a Intuitive understanding of bias-variance.
 - 8.15. Best and worst cases for an algorithm.
 - 8.16. Performance measurement of models:
 - 8.17. Accuracy
 - 8.18. Confusion matrix, TPR, FPR, FNR, TNR
 - 8.19. Precision & recall, F1-score.
 - 8.20. Receiver Operating Characteristic Curve (ROC) curve and AUC.
 - 8.21. Log-loss.
 - 8.22. R-Squared/ Coefficient of determination.
 - 8.23. Median absolute deviation (MAD)
 - 8.24. Distribution of errors.

- 9. Naive Bayes
 - 9.1. Conditional probability.
 - 9.2. Independent vs Mutually exclusive events.
 - 9.3. Bayes Theorem with examples.
 - 9.4. Exercise problems on Bayes Theorem.
 - 9.5. Naive Bayes algorithm.
 - 9.6. Toy example: Train and test stages.
 - 9.7. Naive Bayes on Text data.
 - 9.8. Laplace/Additive Smoothing.
 - 9.9. Log-probabilities for numerical stability.
 - 9.10. Cases:
 - 9.10.1. Bias and Variance tradeoff.
 - 9.10.2. Feature importance and interpretability.
 - 9.10.3. Imbalanced data.
 - 9.10.4. Outliers.
 - 9.10.5. Missing values.
 - 9.10.6. Handling Numerical features (Gaussian NB)
 - 9.10.7. Multiclass classification.
 - 9.10.8. Similarity or Distance matrix.
 - 9.10.9. Large dimensionality.
 - 9.10.10. Best and worst cases.
 - 9.11. Code example
 - 9.12. Exercise: Apply Naive Bayes to Amazon reviews.
- 10. Logistic Regression:
 - 10.1. Geometric intuition.
 - 10.2. Sigmoid function & Squashing
 - 10.3. Optimization problem.
 - 10.4. Weight vector.
 - 10.5. L2 Regularization: Overfitting and Underfitting.
 - 10.6. L1 regularization and sparsity.
 - 10.7. Probabilistic Interpretation: GaussianNaiveBayes
 - 10.8. Loss minimization interpretation
 - 10.9. Hyperparameter search: Grid Search and Random Search
 - 10.10. Column Standardization.
 - 10.11. Feature importance and model interpretability.
 - 10.12. Collinearity of features.
 - 10.13. Train & Run time space and time complexity.
 - 10.14. Real world cases.
 - 10.15. Non-linearly separable data & feature engineering.
 - 10.16. Code sample: Logistic regression, GridSearchCV, RandomSearchCV
 - 10.17. Exercise: Apply Logistic regression to Amazon reviews dataset.
 - 10.18. Extensions to Logistic Regression: Generalized linear models (GLM)

11. Linear Regression and Optimization.
 - 11.1. Geometric intuition.
 - 11.2. Mathematical formulation.
 - 11.3. Cases.
 - 11.4. Code sample.
 - 11.5. Solving optimization problems
 - 11.5.1. Differentiation.
 - 13.5.1_a Online differentiation tools
 - 11.5.2. Maxima and Minima
 - 11.5.3. Vector calculus: Grad
 - 11.5.4. Gradient descent: geometric intuition.
 - 11.5.5. Learning rate.
 - 11.5.6. Gradient descent for linear regression.
 - 11.5.7. SGD algorithm.
 - 11.5.8. Constrained optimization & PCA
 - 11.5.9. Logistic regression formulation revisited.
 - 11.5.10. Why L1 regularization creates sparsity?
 - 11.5.11. Exercise: Implement SGD for linear regression.
12. Support Vector Machines (SVM)
 - 12.1. Geometric intuition.
 - 12.2. Mathematical derivation.
 - 12.3. Loss minimization: Hinge Loss.
 - 12.4. Dual form of SVM formulation.
 - 12.5. Kernel trick.
 - 12.6. Polynomial kernel.
 - 12.7. RBF-Kernel.
 - 12.8. Domain specific Kernels.
 - 12.9. Train and run time complexities.
 - 12.10. nu-SVM: control errors and support vectors.
 - 12.11. SVM Regression.
 - 12.12. Cases.
 - 12.13. Code Sample.
 - 12.14. Exercise: Apply SVM to Amazon reviews dataset.
13. Decision Trees
 - 13.1. Geometric Intuition: Axis parallel hyperplanes.
 - 13.2. Sample Decision tree.
 - 13.3. Building a decision Tree:
 - 13.3.1. Entropy
 - 15.3.1.a Intuition behind entropy
 - 13.3.2. Information Gain
 - 13.3.3. Gini Impurity.

- 13.3.4. Constructing a DT.
 - 13.3.5. Splitting numerical features.
 - 15.3.5a Feature standardization.
 - 13.3.6. Categorical features with many possible values.
- 13.4. Overfitting and Underfitting.
- 13.5. Train and Run time complexity.
- 13.6. Regression using Decision Trees.
- 13.7. Cases
- 13.8. Code Samples.
- 13.9. Exercise: Decision Trees on Amazon reviews dataset.
- 14. Ensemble Models:
 - 14.1. What are ensembles?
 - 14.2. Bootstrapped Aggregation (Bagging)
 - 14.2.1. Intuition
 - 14.2.2. Random Forest and their construction.
 - 14.2.3. Bias-Variance tradeoff.
 - 14.2.4. Train and Run-time Complexity.
 - 14.2.5. Code Sample.
 - 14.2.6. Extremely randomized trees.
 - 14.2.7. Cases
 - 14.3. Boosting:
 - 14.3.1. Intuition
 - 14.3.2. Residuals, Loss functions and gradients.
 - 14.3.3. Gradient Boosting
 - 14.3.4. Regularization by Shrinkage.
 - 14.3.5. Train and Run time complexity.
 - 14.3.6. XGBoost: Boosting + Randomization
 - 14.3.7. AdaBoost: geometric intuition.
 - 14.4. Stacking models.
 - 14.5. Cascading classifiers.
 - 14.6. Kaggle competitions vs Real world.
 - 14.7. Exercise: Apply GBDT and RF to Amazon reviews dataset.
- 15. Featurizations and Feature engineering.**
 - 15.1. Introduction.
 - 15.2. Time-series data.
 - 15.2.1. Moving window.
 - 15.2.2. Fourier decomposition.
 - 15.2.3. Deep learning features: LSTM
 - 15.3. Image data.
 - 15.3.1. Image histogram.
 - 15.3.2. Keypoints: SIFT.

- 15.3.3. Deep learning features: CNN
- 15.4. Relational data.
- 15.5. Graph data.
- 15.6. Feature Engineering.
 - 15.6.1. Indicator variables.
 - 15.6.2. Feature binning.
 - 15.6.3. Interaction variables.
 - 15.6.4. Mathematical transforms.
- 15.7. Model specific featurizations.
- 15.8. Feature orthogonality.
- 15.9. Domain specific featurizations.
- 15.10. Feature slicing.
- 15.11. Kaggle Winners solutions.

17a. Miscellaneous Topics

- 17a.1 Calibration of Models.
 - 17a.1.1 Need for calibration.
 - 17a.1.2 Calibration Plots.
 - 17a.1.3 Platt's Calibration/Scaling.
 - 17a.1.4 Isotonic Regression
 - 17a.1.5 Code Samples
 - 17a.1.6 Exercise: Calibration + Naive Bayes.
- 15.12. Modeling in the presence of outliers: RANSAC
- 15.13. Productionizing models.
- 15.14. Retraining models periodically as needed.
- 15.15. A/B testing.
- 15.16. VC Dimensions.

16. Unsupervised learning/Clustering: K-Means (2)

- 16.1. What is Clustering?
- 16.2. Unsupervised learning
- 16.3. Applications.
- 16.4. Metrics for Clustering.
- 16.5. K-Means
 - 16.5.1. Geometric intuition, Centroids
 - 16.5.2. Mathematical formulation: Objective function
 - 16.5.3. K-Means Algorithm.
 - 16.5.4. How to initialize: K-Means++
 - 16.5.5. Failure cases/Limitations.
 - 16.5.6. K-Medoids
 - 16.5.7. Determining the right K.
 - 16.5.8. Time and space complexity.

- 16.5.9. Code Samples
- 16.5.10. Exercise: Cluster Amazon reviews.

17. Hierarchical clustering

- 17.1. Agglomerative & Divisive, Dendrograms
- 17.2. Agglomerative Clustering.
- 17.3. Proximity methods: Advantages and Limitations.
- 17.4. Time and Space Complexity.
- 17.5. Limitations of Hierarchical Clustering.
- 17.6. Code sample.
- 17.7. Exercise: Amazon food reviews.

18. DBSCAN (Density based clustering)

- 18.1. Density based clustering
- 18.2. MinPts and Eps: Density
- 18.3. Core, Border and Noise points.
- 18.4. Density edge and Density connected points.
- 18.5. DBSCAN Algorithm.
- 18.6. Hyper Parameters: MinPts and Eps.
- 18.7. Advantages and Limitations of DBSCAN.
- 18.8. Time and Space Complexity.
- 18.9. Code samples.
- 18.10. Exercise: Amazon Food reviews.

19. **Recommender Systems and Matrix Factorization. (3)**

- 19.1. Problem formulation: Movie reviews.
- 19.2. Content based vs Collaborative Filtering.
- 19.3. Similarity based Algorithms.
- 19.4. Matrix Factorization:
 - 19.4.1. PCA, SVD
 - 19.4.2. NMF
 - 19.4.3. MF for Collaborative filtering
 - 19.4.4. MF for feature engineering.
 - 19.4.5. Clustering as MF
- 19.5. Hyperparameter tuning.
- 19.6. Matrix Factorization for recommender systems: Netflix Prize Solution
- 19.7. Cold Start problem.
- 19.8. Word Vectors using MF.
- 19.9. Eigen-Faces.
- 19.10. Code example.
- 19.11. Exercise: Word Vectors using Truncated SVD.

20. **Neural Networks.**

- 20.1. History of Neural networks and Deep Learning.
- 20.2. Diagrammatic representation: Logistic Regression and Perceptron
- 20.3. Multi-Layered Perceptron (MLP).
- 20.4. Training an MLP.
- 20.5. Backpropagation.
- 20.6. Weight initialization.
- 20.7. Vanishing Gradient problem.
- 20.8. Bias-Variance tradeoff.
- 20.9. Decision surfaces: Playground
Refer: playground.tensorflow.org
- 20.10. Tensorflow and Keras.
 - 20.10.1. Introduction to tensorflow
 - 20.10.2. Introduction to Keras.
 - 20.10.3. Computational Graph.
 - 20.10.4. Building and MLP from scratch.
 - 20.10.5. MLP in Keras.
 - 20.10.6. GPU vs CPU.
 - 20.10.7. GPUs for Deep Learning.
- 20.11. Assignment: MNIST using MLP.

21. Deep Multi-layer perceptrons

- 21.1. Dropout layers & Regularization.
- 21.2. Rectified Linear Units (ReLU).
- 21.3. Batch Normalization.
- 21.4. Optimizers:
 - 21.4.1. Local and Global Optima.
 - 21.4.2. ADAM
 - 21.4.3. RMSProp
 - 21.4.4. AdaGrad
- 21.5. Gradient Checking.
- 21.6. Initialization of models.
- 21.7. Softmax and Cross-entropy for multi-class classification.
- 21.8. Code Sample: MNIST
- 21.9. Auto Encoders.
- 21.10. Word2Vec.

22. Convolutional Neural Nets.

- 22.1. MLPs on Image data
- 22.2. Convolution operator.
- 22.3. Edge Detection on images.
- 22.4. Convolutional layer.
- 22.5. Max Pooling & Padding.
- 22.6. Imagenet dataset.
- 22.7. AlexNet.
- 22.8. Residual Network.

- 22.9. Inception Network.
- 22.10. Transfer Learning: Reusing existing models.
 - 22.10.1. How to reuse state of the art models for your problem.
 - 22.10.2. Data Augmentation.
 - 22.10.3. Code example: Cats vs Dogs.
- 23. Long Short-term memory (LSTMs)**
 - 23.1. Recurrent Neural Network.
 - 23.2. Backpropagation for RNNs.
 - 23.3. Memory units.
 - 23.4. LSTM.
 - 23.5. GRUs.
 - 23.6. Bidirectional RNN.
 - 23.7. Code example: Predict a stock price using LSTM.
 - 23.8. Sequence to Sequence Models.
- 24. Case Study: Personalized Cancer Diagnosis.**
 - 24.1. Business/Real world problem
 - 24.1.1. Overview.
 - 24.1.2. Business objectives and constraints.
 - 24.2. ML problem formulation
 - 24.2.1. Data
 - 24.2.2. Mapping real world to ML problem.
 - 24.2.3. Train, CV and Test data construction.
 - 24.3. Exploratory Data Analysis
 - 24.3.1. Reading data & preprocessing
 - 24.3.2. Distribution of Class-labels.
 - 24.3.3. "Random" Model.
 - 24.3.4. Univariate Analysis
 - 24.3.4.1. Gene feature.
 - 24.3.4.2. Variation Feature.
 - 24.3.4.3. Text feature.
 - 24.4. Machine Learning Models
 - 24.4.1. Data preparation.
 - 24.4.2. Baseline Model: Naive Bayes
 - 24.4.3. K-Nearest Neighbors Classification.
 - 24.4.4. Logistic Regression with class balancing
 - 24.4.5. Logistic Regression without class balancing
 - 24.4.6. Linear-SVM.
 - 24.4.7. Random-Forest with one-hot encoded features
 - 24.4.8. Random-Forest with response-coded features
 - 24.4.9. Stacking Classifier
 - 24.4.10. Majority Voting classifier.
 - 24.5. Assignments.

25. Taxi demand prediction in New York City.

- 25.1. Business/Real world problem.
 - 25.1.1. Overview
 - 25.1.2. Objectives and Constraints
- 25.2. Mapping to ML problem
 - 25.2.1. Data
 - 25.2.2. dask dataframes
 - 25.2.3. Fields/Features.
 - 25.2.4. Time series forecasting/Regression.
 - 25.2.5. Performance metrics.
- 25.3. Data Cleaning
 - 25.3.1. Latitude and Longitude data
 - 25.3.2. Trip Duration.
 - 25.3.3. Speed.
 - 25.3.4. Distance.
 - 25.3.5. Fare.
 - 25.3.6. Remove all outliers/erroneous points.
- 25.4. Data Preparation
 - 25.4.1. Clustering/Segmentation
 - 25.4.2. Time binning
 - 25.4.3. Smoothing time-series data.
 - 25.4.4. Time series and Fourier transforms.
- 25.5. Baseline models
 - 25.5.1. Ratios and previous-time-bin values.
 - 25.5.2. Simple moving average.
 - 25.5.3. Weighted Moving average.
 - 25.5.4. Exponential weighted moving average.
 - 25.5.5. Results.
- 25.6. Regression models:
 - 25.6.1. Train-Test split & Features
 - 25.6.2. Linear regression.
 - 25.6.3. Random Forest regression.
 - 25.6.4. Xgboost Regression.
 - 25.6.5. Model comparison.
- 25.7. Assignment.

26. Case Study: Microsoft Malware Detection

- 26.1. Business/real world problem
 - 26.1.1. Problem definition
 - 26.1.2. Objectives and constraints.
- 26.2. Machine Learning problem mapping
 - 26.2.1. Data overview.
 - 26.2.2. ML problem.
 - 26.2.3. Train and test splitting.

- 26.3. Exploratory Data Analysis
 - 26.3.1. Class distribution.
 - 26.3.2. Feature extraction from byte files
 - 26.3.3. Multivariate analysis of features from byte files.
 - 26.3.4. Train-Test class distributions
- 26.4. ML models - using byte files only
 - 26.4.1. Random Model.
 - 26.4.2. k-NN
 - 26.4.3. Logistic regression
 - 26.4.4. Random Forest & Xgboost
- 26.5. ASM Files
 - 26.5.1. Feature extraction & Multi-threading.
 - 26.5.2. File-size feature
 - 26.5.3. Univariate analysis on ASM features.
 - 26.5.4. t-SNE analysis.
 - 26.5.5. ML models on ASM file features
- 26.6. Models on all features
 - 26.6.1. t-SNE
 - 26.6.2. RandomForest and Xgboost
- 26.7. Assignments.
- 27. Case studies/Projects:**
 - 27.1. Amazon fashion discovery engine.
 - 27.2. Song Similarity engine.
 - 27.3. Predict customer propensity to purchase using CRM data.
 - 27.4. Suggest me a movie to watch: Netflix Prize.
 - 27.5. Human Activity Recognition using mobile phone's accelerometer and gyroscope data.
 - 27.6. Which ad to show to which user: Ad Click prediction.